
Predicting Building Permit Issue Times

*A Data Science Project
By Aparna Shastry*

Content

- Business Problem
- Data / Data Wrangling
- Exploratory Data Analysis
- Predictive Modeling
- Conclusion / Future Work

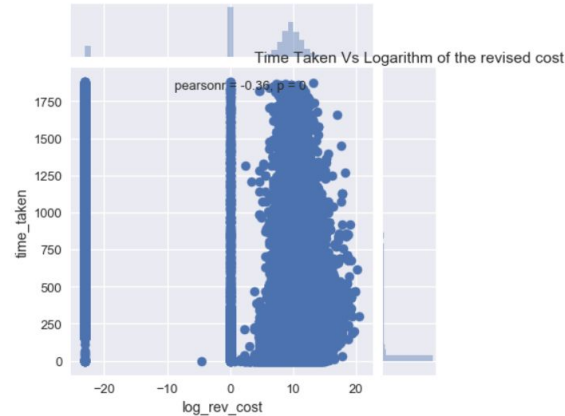

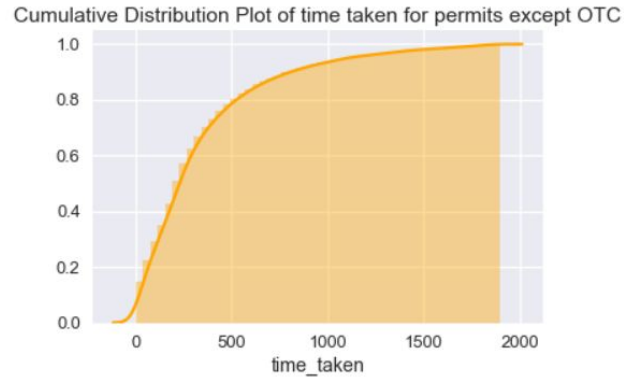
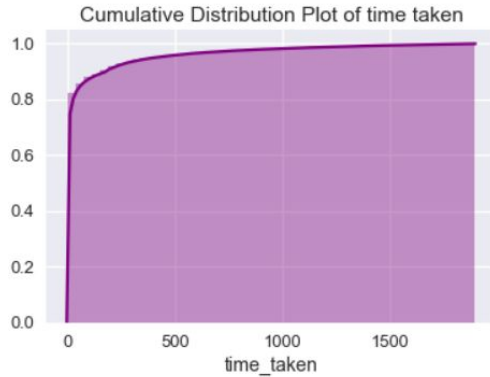
The Business Problem

- Building Permit: A Document issued by city council to builders
- **Delays in issuance:** Major inconvenience and Revenue loss
 - [Trulia Study](#)
 - [Vancouver City Problems](#)
- **Aim of this Study:** Explore the data and model the delays
- **Possible Clients:** Builders / Contractors / Real Estate Agencies
- **How this can help:**
 - Uncertainty Reduction
 - Better planning
 - Proactive follow up by expecting delays

Data / Data Wrangling

- Any building permit application data
- Took [San Francisco city open data](#), for prototyping
- Raw data: **198900k records, 43 columns**
 - Download Date Feb 25th, 2018
 - Most recent 5 years
- **Cleaning:**
 - Retain useful columns and rows
 - Fill a few blank cells with “N”
 - Leave a few blank cells as it is
 - Convert object to dates/float when applicable
 - Change invalid weekdays into valid entries
- After clean up, left with **180811 records, 16 useful features**

Exploratory Data Analysis: Visualizations




Exploratory Data Analysis: Key Findings

- Monday, the least crowded day.
- 90% of permits are OTC alterations.
 - 60%+ issued the same day.
 - 75% of issued within 5 days.
- Median time for new construction type: about 1 year 3 months
- Median of time taken for alterations: close to 6 months
- Most applications have one of the 2 Plansets (although there are 7)
- Absence of cost entry: Major cause of delay
 - ***Fill the cost field.***

Predictive Modeling

- **Target variable Y:**
 - $Y = 0$ if time < 8 days
 - 1 elseif time < 92 days
 - 2 else
- Why not estimate in days (regression)?
 - Too many outliers
 - Records without issue dates
- Why not just have 2 classes?
 - Too trivial, several permit types, not a true representative
- **Metrics used:** Accuracy, Recall and weighted F1 score



Training / Testing / Feature Engineering

- Data split: Train / Validation / Test in the ratio 60:20:20
- Hyperparameter tuned with 5-fold cross validation


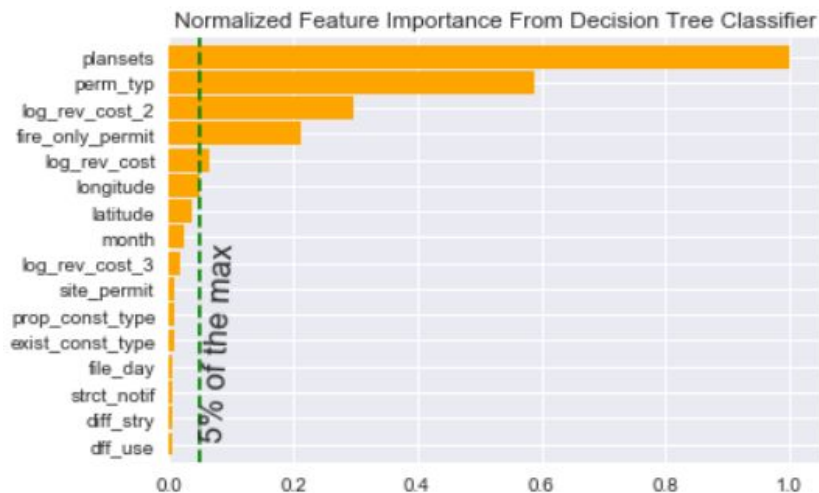
Feature Engineering

- Log of revision cost from revision cost
- Square and cube of log of revision cost
- `dff_use` as `existing_use != proposed_use`,
- `diff_story` as `existing_use != proposed_story`
- Week day and month extracted from filing date
- Location opened up as latitude and longitude

Comparisons of Models Tried

Classifier Model	Hyperparameters	Sensitivity on labels 0,1,2	F1 score
Logistic Regression	C = 0.1, class_Weight = balanced	0.82,0.75,0.67 Overall 79%	0.81
Decision Tree	max_depth=12,min_samples_leaf=6,class_weight = None	0.94,0.41,0.67 Overall 83%	0.82
Bagging	n_estimators=50	0.94,0.51,0.73 Overall 84.6%	0.84
Random Forests	min_samples_leaf=2,class_weight = balanced,n_estimators=200	0.87,0.69,0.73 Overall 83%	0.84
Gradient Boosting	min_samples_leaf=2,n_estimators=50	0.94,0.65,0.67 Overall 82%	0.83

Feature importance



Conclusions / Future work

Conclusions

- **The Random Forest** is the chosen model. Performance on test set was 83% accuracy, no surprises.
- 5 features are the key deterministic factors: Plansets, Permit type, log_rev_cost_2, longitude, latitude
- Rest of them together add around 2% accuracy

Future Work

- Collect housing/crime data to increase features
- Undersampling, oversampling techniques to handle class imbalance

[Code](#) [Detailed Report](#) (**References in Report**)